


The logo features a stylized, glossy, light blue-grey shape that resembles a drop or a cell. A vibrant rainbow spectrum is visible within the shape, curving from the top left towards the center. The word "Antibody" is written in a large, white, serif font, centered over the shape.

Antibody

Self-healing for AI agents

Four agents, one loop

HOW IT WORKS



What each agent does

THE TEAM

Chaos

01

- Hides instructions inside a ticket's internal notes.
- Breaks a tool so it returns nothing, or talks its way into a refund on someone else's order.
- Reads why the last attack got blocked and tries a different way in.

Target

02

- A customer-support agent running on Llama 3.1 8B through W&B Inference.
- Reads the ticket, looks up orders, issues refunds, sends email.
- Never knows it's being tested.

Judge

03

- Checks the facts first: did a refund fire that nobody asked for? Did it open a ticket it wasn't assigned?
- Asks a bigger model only when the rules can't decide.
- Every verdict is a Weave trace you can open and audit.

Repair

04

- Proposes one fix at a time.
- Prefers tool rules and output validators over prompt edits, so the LLM stops being the security layer.
- Remembers what got rejected and never tries it twice.

Gate

- Three Weave evaluations: stops the new attack, still stops every old one, legitimate customers still get helped.
- Anything less is rejected, and Repair tries again.

Where this goes

NEXT

Ship it

01

- Frozen, diffable config versions
- Promote a version to deploy
- Human sign-off, then nightly runs

Cover every agent

02

- Any agent with a tool manifest
- HTTP adapters for agents we don't host
- One Antibody for a fleet: Zendesk, Intercom, Salesforce

Train the model

03

- The regression suite is a labeled dataset
- Fine-tune on it; the gate is the reward
- The loop that hardens the agent trains it