




Antibody

Self-healing for AI agents

Four agents, one loop

HOW IT WORKS



Each pass leaves the agent harder to break and just as helpful as before.

What each agent does

THE TEAM

Chaos

01

- Hides instructions inside a ticket's internal notes.
- Breaks a tool so it returns nothing, or talks its way into a refund on someone else's order.
- Reads why the last attack got blocked and tries a different way in.

Target

02

- A customer-support agent running on Llama 3.1 8B through W&B Inference.
- Reads the ticket, looks up orders, issues refunds, sends email.
- Never knows it's being tested.

Judge

03

- Checks the facts first: did a refund fire that nobody asked for? Did it open a ticket it wasn't assigned?
- Asks a bigger model only when the rules can't decide.
- Every verdict is a Weave trace you can open and audit.

Repair

04

- Proposes one fix at a time.
- Prefers tool rules and output validators over prompt edits, so the LLM stops being the security layer.
- Remembers what got rejected and never tries it twice.

Gate

- Three Weave evaluations: stops the new attack, still stops every old one, legitimate customers still get helped.
- Anything less is rejected, and Repair tries again.

Where this goes

NEXT

Ship it

01

- Frozen, diffable config versions
- Promote a version to deploy
- Human sign-off, then nightly runs

Cover every agent

02

- Any agent with a tool manifest
- HTTP adapters for agents we don't host
- One Antibody for a fleet: Zendesk, Intercom, Salesforce

Train the model

03

- The regression suite is a labeled dataset
- Fine-tune on it; the gate is the reward
- The loop that hardens the agent trains it

Every failure becomes a permanent test. Every patch has to pass all of them. Helpfulness never drops.